

Special Offer | Flat 15% OFF on All Courses | Use Coupon - WHIZSITE15

[Home](#) > [My Courses](#) > [Google Cloud Certified Associate Cloud Engineer](#) > [Practice Test II](#) > [Report](#)

Search Courses



Practice Test II

Completed on 27-January-2021



Attempt

01



Marks Obtained

18 / 50



Your score

36%



Time Taken

00 H 02 M 31 S



Result

Failed

Domains wise Quiz Performance Report

Join us on [Slack community](#)

No	Domain	Total Question	Correct	Incorrect	Unattempted	Marked as Review
1	Setting up a cloud solution environment	6	2	4	0	0
2	Planning and configuring a cloud solution	12	5	7	0	0
3	Deploying and implementing a cloud solution	10	4	6	0	0
4	Ensuring successful operation of a cloud solution	10	3	7	0	0
5	Configure access and security	10	3	7	0	0
6	Other	2	1	1	0	0
Total	All Domain	50	18	32	0	0

Review the Answers

Sorting by All

Question 1

Correct

Domain :Setting up a cloud solution environment

What is the gcloud command to set default zone for compute engine server using gcloud cli?

- A. **gcloud config set compute/zone us-east-1**
- B. **gcloud config configurations set compute/zone us-east-1a**
- ✓ C. **gcloud config set compute/zone us-east1-a** 
- D. **gcloud defaults set compute/zone us-east-1**

Explanation:

Correct Answer - C

The gcloud command to set default zone for compute engine is gcloud config set compute/zone us-east1-a. Hence, C is the correct answer.

Options A, B, and D are incorrect since these are not the valid gcloud commands to set default compute engine zone.

Reference:

<https://cloud.google.com/sdk/gcloud/reference/config/set>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 2

Correct

Domain :Setting up a cloud solution environment

A startup has reached out to you for help with their single page application which they will be running using docker. They want you to suggest the best possible service which can be used by the developers to deploy and manage the application on their own. Which one according to you will be the best suited for their requirement without considering the management of underlying infrastructure?

- A. Cloud Functions
- B. VM Instance
- ✓ C. App Engine 
- D. Google Cloud Storage

Explanation:

Correct Answer - C

C is correct: App Engine's Flexible environment lets you deploy docker based application easily without managing the underlying infrastructure. Hence, it suits best to the client's requirements.

A is incorrect: Cloud Function does not support running docker container.

B is incorrect: VM is also an option to deploy docker applications but it does not fulfil client's requirement of not managing the infrastructure.

D is incorrect: GCS is purely a storage service. You cannot run/install docker applications on it.

Reference:

<https://cloud.google.com/appengine/docs/flexible/>

[Ask our Experts](#)[Rate this Question? !\[\]\(7d1d6890825e83a6a4a51febe2dcc7f3_img.jpg\) !\[\]\(5b78f4d8e2942ab203be44f938cc0a7c_img.jpg\)](#)[View Queries](#)[open ▾](#)**Question 3****Incorrect****Domain :Setting up a cloud solution environment**

You are working for a healthcare industry and for compliance reason your manager has asked you to make sure that all logs are maintained for a minimum of 6 year. Which service can be

used and is supported by Stackdriver Logging?

- A. BigQuery & GCS 
- B. BigTable & Pub/Sub
- ✓ C. CloudSQL & GCS 
- D. Dataflow and BigTable

Explanation:

Answer A

Option A is correct: Stackdriver Logging supports BigQuery and GCS as export option for long term storage.

Option B is incorrect: Only Pub/Sub is supported by Stackdriver.

Option C is incorrect: Only GCS is supported by Stackdriver.

Option D is incorrect: Both the options are not supported by Stackdriver.

Link:

<https://cloud.google.com/logging/docs/export/>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 4

Incorrect

Domain :Setting up a cloud solution environment

Your company has replaced your laptop OS from Windows to Ubuntu and you installed a stable release of Cloud SDK on your machine, however when you run kubectl, you receive an error stating "Command 'kubectl' not found". Which command will you use to install kubectl on Ubuntu?

- A. sudo apt install kubectl
- B. gcloud install kubectl
- ✓ C. sudo yum install kubectl 

- D. gcloud components install kubectl 

Explanation:

Answer D

The correct command to install kubectl is **gcloud components install kubectl**. Hence, the correct answer is D.

Option A, B and C are incorrect: The following commands doesn't support the desired result.

Link:

<https://cloud.google.com/kubernetes-engine/docs/quickstart>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

[open ▾](#)

Question 5

Incorrect

Domain :Setting up a cloud solution environment

Your client wants you to share a screenshot of all the deployments that are/were running on your Google Cloud Deployment Manager. Which command will you use to resolve this?

- A. gcloud deployment-manager deployments list 
- B. gcloud deployment-manager list deployment
- ✓ C. gcloud deployment-manager deployment list 
- D. gcloud deployment-manager list deployments

Explanation:

Answer A

The valid choice to use deployment manager is **gcloud deployment-manager deployments list** command. This will list all the deployments deployed on your Google Cloud Deployment Manager. Hence, Option A is correct.

Option B, C and D are invalid choices or commands

Link:

<https://cloud.google.com/sdk/gcloud/reference/deployment-manager/deployments/list>

Ask our Experts

Rate this Question?  

[View Queries](#)

open ▾

Question 6

Incorrect

Domain :Setting up a cloud solution environment

You have recently joined a startup and have been asked to create a new Google Cloud account for the company along with setting up budget alarm of \$100. You have setup an alarm with 40%, 60%, 90% and 100% threshold. What will happen once the threshold of 100% is reached?

- A. A notification will be sent to Project Owner & Billing Admins.
- B. Your project will deleted because you have depleted all your budget.
- ✓ C. All VMs will be stopped so that you don't exceed your budget. 
- D. A notification will be sent to Billing Admins and Billing Users. 

Explanation:

Answer D

Option D is correct: When billing costs exceed a percentage of your budget, based on the rules you set, alert notifications are sent to billing administrators and billing account users.

Option A is incorrect: Billing alerts are not sent to Project owner.

Option B is incorrect: GCP does will never delete any resources on its own if you clear your dues regularly.

Option C is incorrect: GCP will not stop your VMs or any other resource on crossing the budget.

Link:

<https://cloud.google.com/billing/docs/how-to/budgets>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 7**Correct****Domain :Planning and configuring a cloud solution**

You are working for a genomics company in the US that is planning to conduct data analysis on around 10 TB of data. The company would require 20 vCPUs and 60GB of RAM for the same. You have been asked to figure out the cheapest option to conduct the analysis. Which of the following would you use?

- A. Sustained discount
- B. Committed use discount
- ✓ C. Preemptible instances 
- D. Cloud Functions

Explanation:

Correct Answer - C

C is correct: Preemptible instances are the cheapest option available. They can cut down the cost up to 80% compared to on-demand instances.

A is incorrect: Sustained discounts are given to you only if your VM is running for the complete month.

B is incorrect: Committed instances are cheaper compared to on-demand but expensive when compared to preemptible instances.

D is incorrect: Cloud Function does not support 20vCPUs and 60GB memory, hence will not serve the purpose.

Reference:

<https://cloud.google.com/preemptible-vms/>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 8**Incorrect****Domain :Planning and configuring a cloud solution**

You have MS SQL Server running on a VM with n1-standard-4 machine type. From last few days, you are noticing performance degradation because of more than 95% memory utilization. CPU metrics look absolutely fine and needs no change. Your manager has asked you to increase a small amount of memory keeping cost in mind. What is the most cost-effective approach?

- A. Change machine type to n1-highmem-4
- B. Change machine type to n1-highcpu-10
- ✓ C. Change machine type to n1-standard-6 
- D. Switch to custom machine type and increase memory size 

Explanation:

Correct Answer - D

Option D is correct: Using custom machine type is the best approach as you can just increase the memory as per your requirement.

Option A is incorrect: n1-highmem-4 gives you 26 GB of memory which is more as you require a small amount of increase in memory.

Option B is incorrect: n1-highcpu-10 is not a valid machine type.

Option C is incorrect: n1-standard-6 is not a valid machine type.

Reference:

<https://cloud.google.com/compute/docs/machine-types>

[Ask our Experts](#)**Rate this Question?**  [View Queries](#)[open ▾](#)**Question 9****Correct****Domain :Planning and configuring a cloud solution**

You are hired by an oil company that wants you to lead the migration of Oracle DB and DB2 to

Google Cloud. Which of the following is the best option?

- A. CloudSQL for Oracle and VM for DB2
- B. CloudSQL for both Oracle and DB2
- ✓ C. VM for both Oracle and DB2 
- D. GAE for both Oracle and DB2

Explanation:

Correct Answer - C

Option C is correct: As there is no managed service for both Oracle and DB2, the only option is to install both of them on a VM.

Option A is incorrect: CloudSQL only supports MySQL, SQL Server and Postgres. Therefore, it does not fit our requirement.

Option B is incorrect: CloudSQL neither supports Oracle, nor DB2. Hence, cannot be used to fulfill our requirement.

Option D is incorrect: Google App Engine is only used to deploy applications.

Reference:

<https://cloud.google.com/sql/>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 10

Incorrect

Domain :Planning and configuring a cloud solution

You are working along with a team of developers who are building an accounting application for a startup. Until now, they were using VM disk to store logs but have now realized that this strategy is not scalable and want a different solution to store logs. Which service would you recommend to the team?

- A. Activity logs
- B. Stackdriver logging 

- ✓ C. Flow logs 

- D. GKE

Explanation:

Correct Answer - B

Option B is correct: Stackdriver logging is a native Google Cloud solution for storing logs. The only drawback is that it stores data only for 30 days.

Option A is incorrect: You cannot write or push anything to activity logs. They store each and every activity going on in your GCP account.

Option C is incorrect: Flow logs are used to store network logs and are attached to subnets.

Option D is incorrect: GKE is a service to deploy containerized applications on the cloud.

Reference:

<https://cloud.google.com/logging/>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 11

Correct

Domain :Planning and configuring a cloud solution

You are building up an application for your client which has API, Web, App and DB layer. The client has asked you to decouple the application so that each request is successfully acknowledged and processed by application layer. Which service will you use for the same?

- A. App Engine
- B. Cloud Function
- ✓ C. Pub/Sub 
- D. Dataproc

Explanation:

Answer C

Option C is correct: You will be using Pub/Sub to store the messages in a queue and let the application server pull messages from the queue and process them. In this way you can ensure each and every message/request is processed.

Option A is incorrect: App Engine's only purpose is to deploy application without managing the servers.

Option B is incorrect: Cloud Function is a compute service which lets you run your application without managing the infrastructure.

Option D is incorrect: Dataproc lets you launch managed Hadoop clusters for processing and analysing huge amount of data.

Link:

<https://cloud.google.com/pubsub/>

Ask our Experts

Rate this Question?  

[View Queries](#)

open ▾

Question 12

Incorrect

Domain :Planning and configuring a cloud solution

For last few months your company was testing an ERP application built for one of its client. The product is ready to move to production but before handing over the application your manager has asked you to enable continuous backup for MySQL DB running on VM. How would you do it?

- A. Write a cron job to backup data every minute.
- B. Migrate to CloudSQL as there is no other option.
- C. Create a clone of MySQL VM and write shell script to handle the replication/backup. 
- D. Binary logging. 

Explanation:

Answer D

Option D is correct: The binary log contains "events" that describe database changes such as

table creation operations or changes to table data.

Option A is incorrect: Writing a cron job is inefficient way as it will hamper user performance.

Option B is incorrect: As your manager has asked you to enable continuous backup on VM, moving to CloudSQL is not an option.

Option C is incorrect: Running a shell script will affect user performance because you will need to run it continuously for exact replica.

Link:

<https://cloud.google.com/sql/docs/mysql/backup-recovery/restore>

<https://dev.mysql.com/doc/refman/8.0/en/binary-log.html>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 13

Incorrect

Domain :Planning and configuring a cloud solution

A new client has approached your company who wants a serverless application running on Google Cloud. They want their frontend in React and backend to be written in Python because they have in-house team which is familiar with React and Python. Which service(s) will you prefer to use?

- A. Pub/Sub & GCS
- B. Cloud Function & GCS 
- ✓ C. Cloud Function 
- D. GKE & Cloud Function

Explanation:

Answer B

Option B is correct: You can use GCS for hosting a static website (html, css and js) and Cloud Function for deploying Python code.

Option A is incorrect: GCS can be used for hosting serverless frontend application, but

Pub/Sub cannot be used to deploy Python backend. Also, Pub/Sub is used for decoupling the different modules of application.

Option C is incorrect: Using Cloud Function you can only deploy Python code.

Option D is incorrect: Cloud Function can be used to deploy serverless application but GKE will create VMs, hence is not considered for deploying serverless application.

Link:

<https://cloud.google.com/storage/docs/hosting-static-website>

<https://cloud.google.com/functions/>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 14

Correct

Domain :Planning and configuring a cloud solution

Your company is hosting 10TB of customer data in BigQuery. The CTO of company has decided to use this data and build some analytics on top of data which they have. At the end of first month there was a huge spike in the bill due to use of BigQuery and CFO was not happy with the same. He has asked you to cut down the cost. How can you achieve this?

- A. Use GROUP BY clause.
- B. Use composite keys to query the data.
- ✓ C. Instead of using SELECT *, query only required columns. 
- D. Use JOINS in the query to fetch data.

Explanation:

Answer C

Option C is correct: By limiting the number of columns to required ones you will be able to reduce the amount of data processed by BigQuery thus saving cost.

Option A is incorrect: Grouping query result does not help in reducing cost. It just groups result depending on the column specified.

Option B is incorrect: Composite key will not help in reducing cost. This is only used to create

unique records in a table.

Option D is incorrect: JOINS are used when you are querying data from more than one table.

Link:

<https://cloud.google.com/bigquery/pricing>

Ask our Experts

Rate this Question?  

View Queries

[open](#) ▾

Question 15

Incorrect

Domain :Planning and configuring a cloud solution

You created a Nearline bucket in GCS by using the default parameters, after an hour you realised the requirement was to create a Nearline bucket in asia-south1 location. How will you correct this configuration change?

- A. Edit the bucket and change the location of the bucket from default location to asia-south1.
- B. You cannot edit the bucket once it is created
- ✓ C. Modify the bucket and change its storage class to regional and choose the location to asia-south1 
- D. Delete the bucket and create a new bucket as per the requirement. 

Explanation:

Answer D

D is correct: Once the bucket is created, you have no option to change the location of the bucket. The only option is to delete the bucket and create a new bucket as per the requirement.

A is incorrect: This option is not available, and the location is greyed out once the bucket is created

B is incorrect: One can edit a bucket even after it is created.

C is incorrect: This option is invalid as the requirement is to create a bucket in Nearline class.

Link:

<https://cloud.google.com/storage/docs/locations>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 16

Incorrect

Domain :Planning and configuring a cloud solution

You created a CloudSQL instance with automated backup switched on. After few weeks your manager asked you to restore a 12 days older snapshot. While restoring the snapshot you notice that there are only 7 automated backups available. What can be the reason?

- A. You forgot to change the default retention period from 7 days.
- B. CloudSQL only stores 7 automated backups. 
- ✓ C. Looking at the audit logs you find that the snapshot was deleted by one of your team member. 
- D. None of the above.

Explanation:

Answer B

Option B is correct: CloudSQL only stores last 7 automated backups and this setting cannot be changed as of writing this question.

Option A is incorrect: You cannot change the default retention period.

Option C is incorrect: This could have been the correct answer in case you were trying to restore a snapshot which was less than 7 days old.

Option D is incorrect: As the correct answer is B, this is an invalid option.

Link:

<https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 17**Correct****Domain :Planning and configuring a cloud solution**

A client of yours want you to migrate its on-premise MySQL data to Google Cloud without any downtime. Which service will you use for migrating SQL data to Cloud?

- A. **Cloud Migration**
- B. **Anthos**
- ✓ C. **CloudSQL** 
- D. **Cloud Run**

Explanation:

Answer C

Option C is correct: CloudSQL provides database migration service as well as MySQL managed DB service.

Option A is incorrect: There is no such service provided by GCP.

Option B is incorrect: Anthos is used to migrate Kubernetes workloads.

Option D is incorrect: Cloud Run is used to run stateless HTTP containers on a fully managed environment or in your own GKE cluster.

Link:

<https://cloud.google.com/db-migration/>

[Ask our Experts](#)[Rate this Question? !\[\]\(4c035d133c696276f95b45b6582efdbb_img.jpg\) !\[\]\(117237df2fc61235498f32466c67b39c_img.jpg\)](#)[View Queries](#)[open !\[\]\(6d82067ba813a23a033bfcf9dd3cfe39_img.jpg\)](#)**Question 18****Incorrect****Domain :Planning and configuring a cloud solution**

You are starting to work on a client's project who is looking for a database service within Google Cloud which is horizontally scalable, supports gigabyte size of relational data and even supports ACID for reliable storage of data. Which service will you recommend?

- A. Datastore
- B. BigQuery
- ✓ C. CloudSQL 
- D. Cloud Spanner 

Explanation:

Answer D

Option D is correct: Cloud Spanner is a global SQL database service that can scale across the globe horizontally and even supports ACID transactions.

Option A is incorrect: Datastore also supports ACID transaction but is not horizontally scalable and is a NoSQL database.

Option B is incorrect: BigQuery is also a regional service and cannot be scaled horizontally.

Option C is incorrect: CloudSQL supports ACID transactions but supports only vertical scaling.

Link:

<https://cloud.google.com/spanner/>

[Ask our Experts](#)[Rate this Question?](#)  

[View Queries](#)[open](#) 

Question 19**Incorrect**

Domain :Deploying and implementing a cloud solution

One of your team member was conducting an experiment with Kubernetes nodes. He tainted two of six nodes before leaving for the day. Which command will you use to get the list of nodes along with tainted nodes?

- A. **kubectl get nodes** 
- B. **kubectl get nodes --include-tainted**
- ✓ C. **kubectl list nodes --all** 
- D. **kubectl list nodes --include-tainted**

Explanation:

Answer A

The correct command to get the list of nodes including tainted ones is **kubectl get nodes**. Hence, option A is correct.

Option B, C and D are incorrect because they are invalid commands to allow remote access.

Link:

<https://kubernetes.io/docs/tutorials/kubernetes-basics/explore/explore-intro/>

[Ask our Experts](#)[Rate this Question?](#)  [View Queries](#)[open](#) **Question 20****Correct**

Domain :Deploying and implementing a cloud solution

Your program manager wants you to set up a hybrid network between your Google Cloud and on-premise infrastructure that allows high bandwidth, low latency connection between both the networks. As a Sr. Cloud Engineer, which one of the following service will you select:

- A. Virtual Private Network
- B. Cloud Router
- ✓ C. Interconnect 
- D. Shared VPC

Explanation:

Answer C

C is correct: Interconnect lets you establish high bandwidth, low latency connections between your Google Cloud VPC networks and your on-premises infrastructure.

A is incorrect: A virtual private network lets you securely connect your Google Compute Engine resources to your own private network. Google VPN uses IKEv1 or IKEv2 to establish the IPSec connectivity.

B is incorrect: Google Cloud Router enables you to dynamically exchange routes between your Virtual Private Cloud (VPC) and on-premises networks by using Border Gateway Protocol (BGP)

D is incorrect: Shared VPC lets you share subnets with other projects. You can then create resources (like VM instances) on those subnets.

Link:

<https://cloud.google.com/interconnect/docs/>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 21

Incorrect

Domain :Deploying and implementing a cloud solution

Your Cloud Function is using python modules to connect to your SQL PostgreSQL database in

your VPC. As such, how do you bundle your Cloud Function to add the dependencies?

- A. Upload the function and import the dependencies in code so that the Cloud Function can resolve the dependencies automatically.
- B. Upload the code through the Google Cloud console and upload the dependencies as the zip
- C. Zip the function and the dependencies separately and upload them in Cloud Function as multi parts X
- D. Put the function and the dependencies in one folder and zip them together ✓

Explanation:

Answer D

D is correct: Cloud function dependencies could be resolved by using the pip package manager's requirements.txt file or packaging local dependencies alongside your function.

A is incorrect: Cloud Function does not download the dependencies based on the import statement in python

B & C are incorrect: Both the approaches are inappropriate to specify the dependencies of a Cloud Function

Link:

<https://cloud.google.com/functions/docs/concepts/python-runtime>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 22

Incorrect

Domain :Deploying and implementing a cloud solution

Your company is about to migrate few of its server initially from on-premise to Google Cloud and you have been asked to lead this task. All the instances have been migrated completely, but you are still left with creating health checks for all the instances. Which service will help you create health checks?

- A. Stackdriver Monitoring ✓

B. Stackdriver Alerts

- ✓ C. Grafana 
- D. CloudWatch

Explanation:

Answer A

Option A is correct: Stackdriver monitoring helps you monitor GCP resources and create alerts for metrics such as CPU Utilization, Disk Usage, etc.

Option B is incorrect: This is an invalid service.

Option C is incorrect: Grafana is a third-party service and is only used as visualization tool.

Option D is incorrect: CloudWatch service is an AWS service.

Link:

<https://cloud.google.com/monitoring/uptime-checks/>

[Ask our Experts](#)[Rate this Question?](#)  [View Queries](#)[open ▾](#)**Question 23****Correct****Domain :Deploying and implementing a cloud solution**

You have been tasked with developing an application and later hosting it on Google Cloud using a service whose infrastructure scales automatically and is completely managed by Google along with support of Go lang and events. Which service will you use?

- A. App Engine Flexible
- B. Managed VM Instance
- ✓ C. Cloud Functions 
- D. App Engine Standard

Explanation:

The correct answer is Option C.

Option C is correct: Google Cloud Functions is a serverless execution environment for building and connecting cloud services. With Cloud Functions you write simple, single-purpose functions that are attached to events emitted from your cloud infrastructure and services.

Option A is incorrect: Using App Engine Flexible you can host Go lang application, but the underlying infrastructure is not fully managed by Google Cloud.

Option B is incorrect: Even using managed VM instance involves hum intervention.

Option D is incorrect: Using App Engine Standard you can host Go lang application, but the underlying infrastructure is not fully managed by Google Cloud.

Link:

<https://cloud.google.com/functions/docs/concepts/overview>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 24

Incorrect

Domain :Deploying and implementing a cloud solution

Your company is running a tightly coupled application on Google Cloud VM instances, storing each and every data on disks and auto scaling is not possible. Black Friday sale is approaching and your company has provisioned enough number of additional servers and ample amount of disk space to cater the demand but suddenly during the sale you notice that on one of the server disk space is about to hit 90% usage and your manager has asked you to add 200GB immediately.

While upgrading the disk in hurry you accidentally increased the size to 2000GB instead of 200GB and your manager is yelling at you to downgrade the disk. You are trying to downgrade the disk but are unable to do so. What can be the reason?

- A. You need to stop the VM and downgrade the disk.
- B. You can only upgrade the disk. 
- ✓ C. Downgrading the disk is only possible via CLI/API. 
- D. None of the above.

Explanation:

The correct answer is Option B.

Option B is correct: Disks can only be upgraded in terms of size.

Option A is incorrect: As downgrading the disk is not supported stopping the instance is of no help.

Option C is incorrect: Downgrading the disk is not even possible via CLI/API.

Option D is incorrect: Because correct answer is option B, this is an invalid choice.

Link:

<https://cloud.google.com/compute/docs/disks/add-persistent-disk>

Ask our Experts

Rate this Question?  

[View Queries](#)

open 

Question 25

Incorrect

Domain :Deploying and implementing a cloud solution

There are 15 VMs in your company's Google Cloud account which have public IP address but are ephemeral. Your manager asked you to assign static public IP to each instance. During the process you were able to assign IPs to only 5 of the VMs and were receiving error of all other instances. What can be the reason?

- A. You have reached the hard limit of using static IP addresses.
- B. Your quota has been exhausted and you can raise a request to increase it. 
- ✓ C. Not all OS supports static public IP so check if the OS on VM supports static IP. 
- D. You need to first remove existing public IP attached to VMs and then attach static IP to them.

Explanation:

The correct answer is Option B

Option B is correct: Cloud providers have soft limits of some of the resources to prevent uncontrolled usage of resources. Likewise, even GCP has it for static public IPs. You can simply raise request for increasing the limit.

Option A is incorrect: 5 static public IP is not the hard limit because it is very low.

Option C is incorrect: This has nothing to do which OS your VM is hosting.

Option D is incorrect: Removing existing public IP will not help with assigning static public IP to the VM.

Link:

<https://cloud.google.com/compute/quotas>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 26

Incorrect

Domain :Deploying and implementing a cloud solution

Your company has received a new project where it needs to migrate on-premise servers and data to Google Cloud gradually but until then you need to setup a VPN tunnel between on-premise and Google Cloud. Which service will you use in conjunction with Cloud VPN for a smooth setup?

- A. Cloud CDN
- B. Cloud NAT
- ✓ C. Cloud Run 
- D. Cloud Router 

Explanation:

Answer D

Option D is correct: Google Cloud Router enables you to dynamically exchange routes between your Virtual Private Cloud (VPC) and on-premises networks by using Border Gateway Protocol (BGP). The Cloud Router automatically learns new subnets in your VPC network and announces them to your on-premises network.

Option A is incorrect: Cloud CDN leverages Google's globally distributed edge points of presence to accelerate content delivery for websites and applications served out of Google Compute Engine and Google Cloud Storage. Cloud CDN lowers network latency, offloads origins, and reduces serving costs.

Option B is incorrect: Cloud NAT enables you to provision your application instances without public IP addresses while also allowing them to access the internet for updates, patching, config management, and more in a controlled and efficient manner.

Option C is incorrect: Cloud Run is a managed compute platform that automatically scales your stateless containers.

Link:

<https://cloud.google.com/router/docs/>

Ask our Experts

Rate this Question?  

[View Queries](#)

open ▾

Question 27

Correct

Domain :Deploying and implementing a cloud solution

Your company is running a high availability deployment "hello-server" in Kubernetes Engine on port 8080 which needs to be exposed to public internet using a load balancer on port 80. Which of the following commands will help to accomplish this deployment?

- A. `kubectl expose deployment hello-server --type LoadBalancer --port 8080 --target-port 80`
- B. `kubectl run deployment hello-server --type LoadBalancer --port 80 --target-port 8080`
- C. `kubectl expose deployment hello-server --type LoadBalancer --port 80 --target-port 8080` 
- D. `kubectl run deployment hello-server --type LoadBalancer --port 8080 --target-port 80`

Explanation:

Answer C

Option C is correct: You can expose the deployment on port 80 by routing the request to target port 8080 using this command.

Option A is incorrect: The request is getting generated at port 8080 and routed to deployment at port 80 which is not correct.

Option B and D are incorrect: run command does not support the type argument and is used

to create a deployment.

Link:

<https://cloud.google.com/kubernetes-engine/docs/quickstart>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 28

Correct

Domain :Deploying and implementing a cloud solution

Your client is planning to deploy an application on Google App Engine and have a requirement to store 1TB of data in schemaless, strongly consistent, ACID compliant database solution. Which of the following would you consider for this?

- A. Google Cloud SQL
- B. Google Cloud Storage
- ✓ C. Google Cloud Datastore 
- D. Google Cloud Bigtable

Explanation:

Answer C

Option C is correct: Datastore provides a NoSQL document database built for automatic scaling, high performance, and ease of application development. It could be used to store 1TB of data.

Option A is incorrect: Cloud SQL is not schemaless.

Option B is incorrect: Cloud Storage is object storage and doesn't support SQL or NoSQL data.

Option D is incorrect: Cloud Bigtable is schemaless but is not suitable for 1TB of data.

Link:

<https://cloud.google.com/datastore/>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 29**Incorrect****Domain :Ensuring successful operation of a cloud solution**

You created a snapshot schedule and attached it to a disk for automating backup of your VM. Few days later when you were looking at the snapshots you notice that the snapshot size is different for each and every snapshot. What can be the reason?

- A. Snapshots are incremental/point-in-time copies of VM. 
- B. Raise a ticket with the Google support regarding the same.
- ✓ C. It is just a UI bug which is not yet solved by Google. 
- D. Increase the timeout window so that disk is backed up properly.

Explanation:

Answer A

Option A is correct: Snapshots are incremental or point-in-time copies of disk which means that from second snapshot only changed blocks are copied and because of this each snapshot size is different.

Option B is incorrect: Raising a support ticket is of no use.

Option C is incorrect: It is not a UI bug. It shows the actual size of snapshot.

Option D is incorrect: There is no such option present within snapshot scheduler.

Link:

<https://cloud.google.com/compute/docs/disks/create-snapshots>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 30**Correct****Domain :Ensuring successful operation of a cloud solution**

A team of developers has been working on a Java 8 application on their local machine but company has now decided to create Windows VM for team and let them work on VM instead of working on local machine. Additionally, your manager has asked to send all stack traces to a common place. Which service would you use to meet the following requirement?

- A. Stackdriver Logging
- B. Stackdriver Profiler
- ✓ C. Cloud Trace 
- D. Stackdriver Debugger

Explanation:

Answer C

Option C is correct: Cloud Trace

Option A is incorrect: Stackdriver Logging allows you to store, search, analyze, monitor, and alert on log data and events from Google Cloud Platform and Amazon Web Services (AWS).

Option B is incorrect: Poorly performing code increases the latency and cost of applications and web services every day, without anyone knowing or doing anything about it. Stackdriver Profiler changes this by continuously analyzing the performance of CPU or memory-intensive functions executed across an application.

Option D is incorrect: Stackdriver Debugger is a feature of Google Cloud Platform that lets you inspect the state of a running application in real time, without stopping or slowing it down.

Link: <https://cloud.google.com/trace/docs/overview>

[Ask our Experts](#)[Rate this Question?](#)  [View Queries](#)[open ▾](#)**Question 31****Correct****Domain :Ensuring successful operation of a cloud solution**

You launched a compute engine in your google cloud project with a firewall rule which has

only port 22 and ICMP enabled to your organization's IP. You can ping the instance but when you try to connect to it via SSH, you get a "Connection Refused" error. What could cause this issue?

- A. Port 22 is in use by another client machine
- B. Request to SSH from organization IP is not valid
- ✓ C. Instance service sshd is not running 
- D. Host machine is stopped

Explanation:

Answer C

C is correct. The only possible reason why ssh is timed out is due to sshd service running on the host machine has stopped. Hence, the service could be started by adding "service sshd start" as a startup script.

A is incorrect: Multiple clients can access host machine on port 22

B is incorrect: Access to the host machine is allowed only from the organization's IP. Successful ping response means the connection is valid

D is incorrect: Stopped host machine does not respond to ping requests.

Link:

<https://cloud.google.com/compute/docs/instances/connecting-to-instance>

[Ask our Experts](#)[Rate this Question?](#)  

[View Queries](#)[open ▾](#)

Question 32**Incorrect**

Domain :Ensuring successful operation of a cloud solution

Your project manager wants you to maintain persistent backup of your app compute instance on a daily basis with auto delete in 7 days, to make it hassle free you use google console for scheduling backups. The backup schedule "daily-backup" was created successfully but backup was not generated automatically. Suggest the step that you missed performing after creating the backup schedule?

- A. Create a gcloud script in the VM and point it to backup schedule daily-backup create above 
- B. Add snapshot schedule daily-backup on the app disk 
- ✓ C. Select the existing snapshot and add daily-backup schedule to auto generate the backup everyday 
- D. Edit VM instance, go to snapshot schedule and select daily-backup to enable backup for that instance.

Explanation:

Answer B

B is correct: The snapshot schedules can only be added to disk. Hence you need to add snapshot schedule name "daily-backup" to the disk that is being used by app compute instance

A is incorrect: This option is not valid

C is incorrect: Snapshot schedule cannot be added in existing snapshot

D is incorrect: Snapshot schedule can only be added in disks and not in vm instance.

Link:

<https://cloud.google.com/compute/docs/disks/scheduled-snapshots>

Ask our Experts

Rate this Question?  

[View Queries](#)

open 

Question 33

Incorrect

Domain :Ensuring successful operation of a cloud solution

You are trying to get Google Kubernetes server configuration of a project by running the command gcloud container get-server-config, when you hit enter, you get an error message: "There was a problem refreshing your current auth tokens: invalid_grant: Bad Request". What would you do to successfully run this command and get the server configuration from Kubernetes?

- A. gcloud auth login 

- B. gcloud config set account [ACCOUNT]
- ✓ C. gcloud config set project [PROJECT_ID] 
- D. gcloud components update

Explanation:

Answer A

The error is produced when the user's gcloud sdk access has expired.

A is correct: gcloud auth login allows you to obtain new google cloud credentials using your existing email id and password to activate your gcloud sdk. Once the login is successful, you can run gcloud container get-server-config

B is incorrect: If the account token has been expired then you cannot set the account id to access the gcloud SDK. You need to login and set up new token.

C is incorrect: If the account token has been expired then you cannot set the project id to access the gcloud SDK. You need to login and set up new token.

D is incorrect: This command is used to update the SDK to its latest version and is invalid to this question.

Link:

<https://cloud.google.com/sdk/gcloud/reference/auth/login>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 34

Incorrect

Domain :Ensuring successful operation of a cloud solution

You have an application running on Google Cloud VM written in C#.Net from last few weeks. Suddenly you experience elevated latency and poor performance among application endpoints. Which service can help you troubleshoot this issue?

- A. Stackdriver Profiler
- B. Stackdriver Monitoring

- ✓ C. Activity Logs X
 - D. Stackdriver Trace ✓
-

Explanation:

Answer D

Option D is correct: Stackdriver Trace is a distributed tracing system that collects latency data from your applications and displays it in the Google Cloud Platform Console. You can track how requests propagate through your application and receive detailed near real-time performance insights.

Option A is incorrect: Profiler help you continuously monitor CPU and heap profiling.

Option B is incorrect: Stackdriver Monitoring provides visibility into the performance, uptime, and overall health of cloud-powered applications.

Option C is incorrect: Activity logs records each and every activity performed by a user or resources like creating, deleting a VM.

Link:

<https://cloud.google.com/trace/>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 35

Incorrect

Domain :Ensuring successful operation of a cloud solution

Your company was bidding on a big data project form last few months and they have finally received the project. The project wants you to deploy Apache Spark clusters on Google Cloud. Which service will you use?

- A. DataFlow
- B. DataProc 
- ✓ C. BigTable 
- D. Cloud Composer

Explanation:

Answer B

Option B is correct: Cloud Dataproc is a fast, easy-to-use, fully managed cloud service for running Apache Spark and Apache Hadoop clusters in a simpler, more cost-efficient way.

Option A is incorrect: Cloud Dataflow is a fully-managed service for transforming and enriching data in stream (real time) and batch (historical) modes with equal reliability and expressiveness.

Option C is incorrect: A petabyte-scale, fully managed NoSQL database service for large analytical and operational workloads. It supports the open source industry standard HBase API.

Option D is incorrect: Cloud Composer is a fully managed workflow orchestration service that empowers you to author, schedule, and monitor pipelines that span across clouds and on-premises data centers. It is built on the popular Apache Airflow open source project.

Link:

<https://cloud.google.com/dataproc/>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 36

Correct

Domain :Ensuring successful operation of a cloud solution

You are working on a project whose development phase has been completed and manager has asked you to create a new project for UAT environment which should be clone of development environment. Dev environment consists of 5 VMs which needs to be replicated to UAT environment. How can you achieve this?

- A. Launch 5 VMs from scratch and install all the libraries present on dev VMs.
- B. Select an instance, click on Actions and use migrate option to launch VM in another project.
- C. Create custom image of each instance in dev environment and while launching VMs in UAT environment under boot disk select Custom images, select your project and choose the relevant image. 
- D. None of the above.

Explanation:

Answer C

Option C is correct: Custom images can be shared among projects hence this is an easy way to replicate VM in another project.

Option A is incorrect: This can be done but is a tedious task to perform if you installed tons of custom libraries on each VM.

Option B is incorrect: There is no such option to migrate VM to another project.

Option D is incorrect: Because option C is correct this is an invalid choice.

Link:

<https://cloud.google.com/compute/docs/images/sharing-images-across-projects>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 37

Incorrect

Domain :Ensuring successful operation of a cloud solution

Suggest the maximum number of containers you can launch within a GKE cluster:

- A. 300,000 

- B. 500,000
- ✓ C. 100,000 
- D. Unlimited

Explanation:

Answer A

Option A is correct: The max limit is 300,000 containers per GKE cluster as mentioned by Google Cloud.

Option B, C and D are incorrect choices as the values are not defined.

Link:

<https://cloud.google.com/kubernetes-engine/quotas>

Ask our Experts

Rate this Question?  

[View Queries](#)

open ▾

Question 38

Incorrect

Domain :Ensuring successful operation of a cloud solution

You are a Jr. Cloud Engineer at John Smith Co. and you were informed about an error "failed to negotiate an api version" while running kubectl command on GCP shell. Which command would you suggest to resolve this issue initially?

- A. gcloud auth application-default login 
- B. gcloud container clusters get-credentials [CLUSTER_NAME]
- ✓ C. sudo gcloud components install kubectl 
- D. gcloud container clusters describe cluster-name

Explanation:

Answer A

Option A is correct: Initially you will check if kubectl has authentication credentials. To do so,

you will run the following command

Option B is incorrect: The following command could be used to mitigate "Connection Refused" error while running kubectl command.

Option C is incorrect: install command is required when you wish to install a google cloud service. In this case kubectl is already installed.

Option D is incorrect: This is an invalid command to this context

Link:

<https://cloud.google.com/kubernetes-engine/docs/troubleshooting>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 39

Correct

Domain :Configure access and security

You are appointed as a head of infrastructure security and the first thing you are asked is to create few IAM users following best security practices. How would you ensure that these users are only able to launch a new instance and list them?

- A. Create Compute Editor role and attach to the users.
- B. Attach Project Editor role to users.
- C. Create custom role with specific permissions and attach it to users. 
- D. Create a service account with required permissions and attach it to users.

Explanation:

Answer C

Option C is correct: You should always follow least privilege principle and assign permission to any resource. Keeping this in mind it is recommended to create custom role and attach it to users.

Option A is incorrect: Compute Editor role will provide many extra permissions other than just launching and list instances.

Option B is incorrect: Project Editor role will allow users to access all the resources with a

project and this is not best security practise.

Option D is incorrect: You cannot attach a service account to user.

Link:

<https://cloud.google.com/sdk/gcloud/reference/app/deploy>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 40

Incorrect

Domain :Configure access and security

Your company is going for an audit and the auditor has approached you to provide them with all the audit logs related to whizlabs-prod project. Which service will you use to extract the logs for auditors?

- A. Flow logs
- B. Activity logs
- ✓ C. Stackdriver logging 
- D. Audit logs 

Explanation:

Answer D

Option A is incorrect: Flow logs provide information about traffic flowing with the VPC.

Option B is incorrect: Activity logs is now deprecated in GCP.

Option C is incorrect: Using only the Stackdriver logging we can not get the activity/user logs which can be share with auditors.

Option D is correct: Under Stackdriver logging, there is a module "audit log". It essentially helps you determine who did what, where, and when. Since the Activity log service is now deprecated, "Audit Log" is the correct answer. We will update this to Option D as the correct answer.

Link:

<https://cloud.google.com/logging/docs/audit/>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 41

Correct

Domain :Configure access and security

You want to filter logs present in Stackdriver Log Viewer using more than simple text filter. Which feature will allow you to do it?

- A. Config Filter
- B. Multi-Stack Filter
- ✓ C. Advanced Filter 
- D. Deep Log Filter

Explanation:

Answer C

Option C is correct: Advance filter lets you filter logs with parameters like severity, resource type, request status, etc.

Option A, B and D are incorrect as there no such options present within Stackdriver Logging for filtering.

Link:

<https://cloud.google.com/logging/docs/view/advanced-filters>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 42

Incorrect

Domain :Configure access and security

You are creating a firewall rule for your DB server so that only web servers and your internal application servers can access it. You have added both service account and tags for traffic filter but while saving you receive an error. What can be the reason?

- A. Service accounts are not supported as filter.
- B. Tags are not supported as traffic filter.
- ✓ C. You can only specify only one traffic filter at a time. 
- D. Service account and tags cannot be used together at the same time. 

Explanation:

Answer D

Option D is correct: You can use the combination of IP range and tags or IP range and service account for traffic filter, but service account and tag combination is not supported at the time of writing this question.

Option A is incorrect: Service accounts are supported as traffic filter.

Option B is incorrect: Even tags are supported for filtering traffic reaching resources.

Option C is incorrect: Google Cloud does supports multiple filter types for firewall provided they are in correct combination.

Link:

https://cloud.google.com/vpc/docs/firewalls#sources_or_destinations_for_the_rule

Ask our Experts

Rate this Question?  

[View Queries](#)

open 

Question 43

Incorrect

Domain :Configure access and security

You have been asked to write a policy in JSON format by your manager. Which one out of below four is the correct way of creating a policy?

- A.
- ```
{
 "bindings": [
 {"roles": [
 "roles/owner"
],
 "members": [
 "user:jon@whizlabs.com",
 "group:hr@whizlabs.com"
]
 }
}
```
- 
- B.
- ```
{  
  "bindings": [  
    {  
      "role": "roles/owner",  
      "members": [  
        "user:jon@whizlabs.com",  
        "group:hr@whizlabs.com"  
      ]  
    }  
  ]  
}
```
- C.
- ```
{
 "bindings": [
 {
 "roles": [
 "roles/owner"
],
 "members": [
 "user:jon@whizlabs.com",
 "group:hr@whizlabs.com"
]
 }
]
}
```
- 
- D.
- ```
{  
  "bindings": [  
    {"role": "roles/owner",  
     "members": [  
       "user:jon@whizlabs.com",  
       "group:hr@whizlabs.com"  
     ]  
  ]  
}
```

Explanation:

Answer B

Option B is the correct way to write a policy for assigning permission.

Option A, C and D are incorrect because they are invalid policies nad have syntax errors.

Link:

<https://cloud.google.com/iam/reference/rest/v1/Policy>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 44

Incorrect

Domain :Configure access and security

Jacob & co has set up a continuous deployment for App Engine using Cloud Build. The builds have worked for over 9 months now but you are approaching the limits of App Engine in terms of how many versions can be stored in the service. What is the maximum number of versions that you can deploy on your app's pricing?

- A. Free App: 15 & Paid App: 210 
- B. Free App: 5 & Paid App: 105
- ✓ C. Free App: 15 & Paid App: 120 
- D. Free App: 5 & Paid App: 150

Explanation:

Answer A

A is correct: As per google docs, Maximum version per app is 15 for free app and 210 for paid app

B is incorrect: The following limit is defined for services in app engine

C & D are incorrect: Both the limits are not defined and invalid

Link:

<https://cloud.google.com/appengine/docs/standard/php/an-overview-of-app-engine#limits>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 45**Correct**

Domain :Configure access and security

Under what scenario will you prefer using tag or service account with firewall instead of using IP address?

- A. Tag or service account can prevent from network sniffing attacks.
- B. Tag or service account are much safer compared to using IP address.
- ✓ C. Any instance launched using tag or service account will have firewall attached automatically. 
- D. None of the above.

Explanation:

Answer C

Option C is correct: Using tag or service account will keep your firewall rule tidy as easy to understand because if you have to apply the same rule to hundreds of instances imagine going through each and every instance and adding their IP in the source. Instead, you can just use common tag or service account among VMs.

Option A is incorrect: This is an invalid choice because it does not prevent from network sniffing attack.

Option B is incorrect: Using IP address instead will be as safer as using tag or service account.

Option D is incorrect: As correct answer is option C, this is an invalid choice.

Link:

<https://cloud.google.com/monitoring/uptime-checks/>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 46**Incorrect**

Domain :Configure access and security

Steve & sons has appointed a new devops engineer with experience in Google Cloud Platform, you have been asked to create a user which would allow him to call "builds.cancel()" API only. Which of the following permission would you consider completing this task?

- A. cloudbuild.builds.create
- B. cloudbuild.builds.update 
- ✓ C. cloudbuild.builds.get 
- D. cloudbuild.builds.list

Explanation:

Answer B

Option B is correct: cloudbuild.builds.update allows user to cancel builds during execution. The following permission is a part of primitive primitive role Cloud Build Editor hence you need to create a custom user and add the specific permission to this.

Option A is incorrect: The following permission allows the user to create, patch, delete the builds

Option C is incorrect: The following permission allows the user to get the builds and triggers only

Option D is incorrect: The following permission allows the user to list the builds and triggers only

Link:

https://cloud.google.com/iam/docs/understanding-roles#primitive_roles

[Ask our Experts](#)Rate this Question?  [View Queries](#)[open ▾](#)**Question 47****Incorrect**

Domain :Configure access and security

John is assigned a new devops project in GCP and he needs deploymentmanager.deployments.* permission to make use to Deployment Manager service. Which of the following role would you create for him?

- A. Deployment Manager Editor
- B. Deployment Manager Custom 
- ✓ C. Deployment Manager Type Viewer 
- D. Deployment Manager Viewer

Explanation:

Answer B

Option B is correct: You would need to create a custom role and add deploymentmanager.deployments.* permissions only

Option A is incorrect: The following role will give full access to deployment manager service

Option C is incorrect: The following role will give Read-only access to all Type Registry resources.

Option D is incorrect: The following role will give Read-only access to all Deployment Manager resources.

Link:

https://cloud.google.com/iam/docs/understanding-roles#primitive_roles

Ask our Experts

Rate this Question?  

[View Queries](#)

open ▾

Question 48

Incorrect

Domain :Configure access and security

Your project manager wants to create a user for Aston Smith who is the new cloud SQL administrator in your organization. Which of the following role would give him the ability to manage specific instances but no ability to import or restore data from backups?

- A. Cloud SQL Editor 

- B. Cloud SQL Admin
- ✓ C. Cloud SQL Viewer ✗
- D. Cloud SQL Client

Explanation:

Answer A

Option A is correct: SQL Editor allows users to Manage specific instances. No ability to see or modify permissions, nor modify users or SSL certs. No ability to import data or restore from a backup, nor clone, delete, or promote instances.

Option B is incorrect: Full control for all Cloud SQL resources.

Option C is incorrect: Read-only access to all Cloud SQL resources.

Option D is incorrect: Connectivity access to Cloud SQL instances from App Engine and the Cloud SQL Proxy. Not required for accessing an instance using IP addresses.

Link:

<https://cloud.google.com/sql/docs/mysql/project-access-control>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 49

Incorrect

Domain : Other

You work for one of the biggest digital media company in USA .The company management has decided to move 90 TB of backups and archival data to Google Cloud. They are looking for long term cost effective archival storage for disaster recovery in Google Cloud . Please select the right solution.

- A. Storage Transfer and Nearline storage
- B. Transfer Appliance and Coldline storage. 
- ✓ C. gsutil and Cloud storage 
- D. Transfer Appliance and Nearline storage.

Explanation:

Answer B

Option B is the correct choice because ,Transfer Appliance is the best choice moving large volume of data and since they are looking for long term cost effective disaster recovery solution , coldline is the best option.

Option A is Incorrect because Storage Transfer is used to import online data into Cloud Storage . Your online data source can be an Amazon Simple Storage Service (Amazon S3) bucket, an HTTP/HTTPS location, or a Cloud Storage bucket

Option C is Incorrect because , gsutil isn't recommended for large volume of data transfer ,It will take a very long time for data transfer depending on the bandwidth.

Option D is Incorrect because , Coldline is a more cost effective archival storage for disaster recovery.

<https://cloud.google.com/transfer-appliance/>

<https://cloud.google.com/storage-transfer/docs/overview>

[Ask our Experts](#)[Rate this Question?](#)  

[View Queries](#)[open ▾](#)

Question 50**Correct**

Domain : Other

You have about 20TB of data which is not accessed and the data team decided to archive them in the cloud. The team is looking for a storage solution that is highly available with minimum costs. On the other hand, the data may be accessed a couple of times a year for reconciliation purposes.

Which of the following choices best satisfy data team's requirements?

- A. Google Storage Standard.
- B. Google Storage Nearline.
- ✓ C. Google Storage Coldline. 
- D. BigQuery.

Explanation:

Answer: C.

Google Cloud Storage Coldline is a low-cost, highly durable storage service for storing infrequently accessed data. Coldline Storage is ideal for data you plan to read or modify *on average once a quarter or less*.

Answer A is incorrect: Google Storage Standard is not a cheap option for storing archive data. There are other options to store archive data cheaper.

Answer B is incorrect: Scenario mentions the need to access the data a couple of times a year. Hence, Nearline is not a suitable option because Nearline is to plan archive data accessed at most once a month.

Answer D is incorrect: BigQuery does not have cheaper cost for storing data than Nearline for the first 90 days until it's moved to long-term storage.

Source(s):

Google Storage Classes: <https://cloud.google.com/storage/docs/storage-classes>

Ask our Experts

Rate this Question?  

View Queries

open 

Finish Review

Certification	Company	Support	Join us on Slack!
Cloud Certification	Become Our Instructor	Contact Us	 Join our open Slack community and get your queries answered instantly! Our experts are online to answer your questions!
Java Certification	Support	Help Topics	
PM Certification	Discussions		
Big Data Certification	Blog		Follow us
	Business		  

© Copyright 2021. Whizlabs Software Pvt. Ltd. All Right Reserved.