

Special Offer | Flat 15% OFF on All Courses | Use Coupon - WHIZSITE15

[Home](#) > [My Courses](#) > [Google Cloud Certified Associate Cloud Engineer](#) > [Practice Test 3](#) > [Report](#)

Search Courses



Practice Test 3

Completed on 27-January-2021



Attempt

01



Marks Obtained

1 / 50



Your score

2%



Time Taken

00 H 00 M 23 S



Result

Failed

Domains wise Quiz Performance Report

Join us on [Slack community](#)

No	Domain	Total Question	Correct	Incorrect	Unattempted	Marked as Review
1	Other	9	0	0	9	0
2	Planning and configuring a cloud solution	13	1	0	12	0
3	Setting up a cloud solution environment	6	0	0	6	0
4	Deploying and implementing a cloud solution	8	0	0	8	0
5	Ensuring successful operation of a cloud solution	7	0	0	7	0
6	Configure access and security	7	0	0	7	0
Total	All Domain	50	1	0	49	0

Review the Answers

Sorting by All

Question 1

Unattempted

Domain : Other

You have a system generated log files required to be later uploaded to Google Storage in the data lake. Since the data is only accessed couple of times in a year by the development team for debugging and log analysis. You are looking for a cheaper storage option for log files than the standard class. Which of the following is suitable?

- A. Cloud Storage Nearline
- B. Cloud Storage Coldline 
- C. BigQuery
- D. Filestore

Explanation:

Correct Answer: B

Google Cloud Storage Nearline is a low-cost, highly durable storage service for storing infrequently accessed data. Nearline Storage is ideal for data you plan to read or modify *on average once a month or less*.

Google Cloud Storage Coldline is a very-low-cost, highly durable storage service for data archiving, online backup, and disaster recovery. Coldline Storage is the best choice for data that you plan to access *at most once a quarter*, due to its slightly lower availability, 90-day minimum storage duration, costs for data access, and higher per operation costs.

Option A is incorrect: Scenario mentions the need to access the data few times a year. Hence, Nearline is not a suitable option because Nearline is used to plan archive data accessed at most once a month.

Options C and D are incorrect: Both BigQuery and Filestore's storage rate is two times higher than Cloud Storage Nearline.

Source(s):

Google Storage Classes:

<https://cloud.google.com/storage/docs/storage-classes>

[Ask our Experts](#)[Rate this Question? !\[\]\(f60b7a900783ac3fd531bfd9c111be6d_img.jpg\) !\[\]\(fe5cf1978663f480c504f8fc2019fe62_img.jpg\)](#)

[View Queries](#)[open !\[\]\(235bfe13ebf007ce2eea9e689707fac7_img.jpg\)](#)

Question 2**Unattempted**

Domain : Other

A system is expected to receive over 15,000 content delivery logs every minute from different web & mobile apps. Logs are received in JSON format. Due to logs being generated by different apps, each developed by a different team, logs do not have a fixed structure and may hold different attributes. Which of the following is a recommended storage option?

- A. Cloud SQL
- B. Cloud Spanner
- C. BigTable 
- D. Datastore

Explanation:

Correct Answer: C

Cloud BigTable is a petabyte-scale, fully managed NoSQL database service for large analytical and operational workloads. It provides flexible schema options.

Options A and B are incorrect: Cloud SQL & Spanner are relational database services. They are not recommended for JSON-format log data with a flexible schema.

Option D is incorrect: Datastore can be a potential choice since it's a NoSQL database. However, Datastore is not built for storing huge data volumes as required in this scenario. Datastore is designed for web applications of small scale.

Source(s):

BigTable vs Datastore:

<https://stackoverflow.com/questions/30085326/google-cloud-bigtable-vs-google-cloud-datastore>

[Ask our Experts](#)[Rate this Question?](#)  [View Queries](#)[open ▾](#)**Question 3****Unattempted**

Domain : Other

A stock market company receives real-time updates from different stock prices in the USA.

The company seeks a solution that can use stock price data for real-time analysis. The solution should allow high throughput to allow queries to run and return the required results with minimum latency. The solution should also be scaled out for more performance. Which of the following products is the best solution in this scenario?

- A. **BigTable** 
- B. **BigQuery**
- C. **Datastore**
- D. **Cloud Spanner**

Explanation:

Correct Answer: A

Cloud BigTable is a petabyte-scale, fully managed NoSQL database service for large analytical and operational workloads. Under a typical workload, Cloud BigTable delivers highly predictable performance. When everything is running smoothly, a typical workload can achieve the following performance for each node in the Cloud Bigtable cluster, depending on which type of storage the cluster uses:

Storage Type	Reads	Writes	Scans
SSD	10,000 rows per second @ 6 ms	or 10,000 rows per second @ 6 ms	220 MB/s
HDD	500 rows per second @ 200 ms	or 10,000 rows per second @ 50 ms	180 MB/s

In general, a cluster's performance increases linearly as you add nodes to the cluster. For example, if you create an SSD cluster with 10 nodes, the cluster can support up to 100,000 rows per second for a typical read-only or write-only workload, with 6 ms latency for each read or write operation.

Option B is incorrect: BigQuery doesn't provide the high throughput and low latency competent to Bigtable. Moreover, you are unable to increase BigQuery's performance, opposed to Bigtable which you can add more nodes for linear performance return.

Option C is incorrect: Datastore is not built for storing and reading huge data volumes as required in this scenario. Datastore is designed for web applications of small scale.

Option D is incorrect: Cloud Spanner does not guarantee the same performance and low latency as BigTable.

Source(s):

Understanding BigTable Performance:

<https://cloud.google.com/bigtable/docs/performance>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 4

Unattempted

Domain : Other

You have an on-premise MySQL database that you have been asked to move to Google Cloud. Users should run SQL queries to fetch data from the database. Your solution should be cost-effective and allow increasing read capacities in the future. Which of the following Google Cloud product is the best for this scenario?

- A. Cloud Storage
- B. Cloud Spanner
- C. Cloud SQL 
- D. Cloud Datastore

Explanation:

Correct Answer: C

Cloud SQL is a fully managed database service that makes it easy to set up, maintain, manage, and administer your relational PostgreSQL, MySQL, and SQL Server databases in the cloud.

Option A is incorrect: Google Storage is blob storage. It does not work as an RDMS.

Option B is incorrect: Cloud Spanner is a very sophisticated and expensive approach for this scenario. Cloud SQL is enough to cover the requirements.

Option D is incorrect: Datastore is a schemaless NoSQL database. Migration is from a structured SQL database so Datastore is not a viable choice.

Source(s):

Cloud SQL:

<https://cloud.google.com/sql/>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 5**Unattempted**

Domain : Other

You need to allow traffic from specific virtual machines in 'subnet-a' network access to machines in 'subnet-b' without giving the entirety of subnet-a access. How can you accomplish this?

- A. Create a firewall rule to allow traffic from resources with specific network tags, then assign the machines in subnet-a the same tags. 
- B. Relocate the subnet-a machines to a different subnet and give the new subnet the needed access.
- C. Create a rule to deny all traffic to the entire subnet, then create a second rule with higher priority giving access to tagged VM's in subnet-a.
- D. You can only grant firewall access to an entire subnet and not individual VM's inside.

Explanation:

Correct Answer A

A (Correct answer) - Create a firewall rule to allow traffic from resources with specific network tags, then assign the machines in subnet-a the same tags.

Network tags allow more granular access based on individually tagged instances - Instances by target tags: The firewall rule is applicable only to VMs if they have a matching network tag.

B - Relocate the subnet-a machines to a different subnet and give the new subnet the needed access.

This would give the entire subnet access which is against the requirements: allow traffic from specific virtual machines in 'subnet-a' network access to machines in 'subnet-b' without giving the entirety of subnet-a access.

C - Create a rule to deny all traffic to the entire subnet, then create a second rule with higher priority giving access to tagged VM's in subnet-a.

Creating overlapping rules with higher priority might technically work, but since traffic defaults to denied if no rule is in place, this is unnecessary. Assigning rules and instances by tags is the

best answer.

D - You can only grant firewall access to an entire subnet and not individual VM's inside.

This is not true per answer A

Ask our Experts

Rate this Question?  

[View Queries](#)

[open ▾](#)

Question 6

Unattempted

Domain : Other

You are managing the GCP Account of a client, the client raises a request to attach 9 local SSDs and launch a VM instance in us-east1 Region, as a Cloud Architect what would be your response to the above request?

- A. You can always attach maximum of ten local SSD devices to a VM instance
- B. If a resource is not available, you won't be able to create new resources of that type, even if you still have remaining quota in your region or project and you can attach up to 24 local SSD devices for 9 TB of total local SSD storage space per instance 
- C. Launch the instance first and add the local SSD drives later for optimal performance.
- D. Request changes to quota from the Quotas page in the GCP Console

Explanation:

Answer: B

Option B is the CORRECT because, Each local SSD is 375 GB in size, but you can attach up to 24 local SSD devices for 9 TB of total local SSD storage space per instance. If a resource is not available, you won't be able to create new resources of that type, even if you still have remaining quota in your region or project.

Option A is INCORRECT because you can only attach up to 8 local SSD devices to a VM instance also if the resource isn't available you can't create the resource.

Option C is INCORRECT because you can attach the local SSD devices to an Instance while launching a VM instance.

Option D is INCORRECT because you can only attach up to 8 local SSD devices to a VM

instance that is the limit.

Read more about it here: <https://cloud.google.com/compute/docs/disks/local-ssd>

<https://cloud.google.com/compute/quotas>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 7

Unattempted

Domain : Other

You have a Kubernetes cluster with 1 node-pool. The cluster receives a lot of traffic and needs to grow. You decide to add a node. What should you do?

- A. Use "gcloud container clusters resize" with the desired number of nodes. 
- B. Use "kubectl container clusters resize" with the desired number of nodes.
- C. Edit the managed instance group of the cluster and increase the number of VMs by 1.
- D. Edit the managed instance group of the cluster and enable autoscaling.

Explanation:

Correct answer A

Feedback

A is correct because this resizes the cluster to the desired number of nodes.

B is not correct because you need to use gcloud, not kubectl.

C is not correct because you should not manually manage the MIG behind a cluster.

D is not correct because you should not manually manage the MIG behind a cluster.

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 8**Unattempted**

Domain : Other

Using the principle of least privilege, your colleague Bob needs to be able to create new instances on Compute Engine in project 'Project A'. How should you give him access without giving more permissions than is necessary?

- A. Give Bob Compute Engine Instance Admin Role for Project A. 
- B. Give Bob Compute Engine Admin Role for Project A.
- C. Create a shared VPC that Bob can access Compute resources from.
- D. Give Bob Project Editor IAM role for Project A.

Explanation:**Correct answer A**

A (Correct answer) - Give Bob Compute Engine Instance Admin Role for Project A - The Compute Engine Instance Admin role only gives access to create/edit compute engine resources such as instances, disks, and snapshots.

B - Give Bob Compute Engine Admin Role for Project A - The Compute Engine Admin role also gives admin access to network resources as well, which is too much.

C - Create a shared VPC that Bob can access Compute resources from – shared VPC is not for granting user role. For example, even if you create shared VPC between "Project A" and "Project B", to meet the requirement, you still need to grant Bob Compute Engine Instance Admin Role for Project A.

D - Give Bob Project Editor IAM role for Project A – that'll be too much for create new instances on Compute Engine in project 'Project A'

Ask our ExpertsRate this Question?  

View Queries

open ▾

Question 9**Unattempted**

Domain : Other

You have an application server running on Compute Engine in the europe-west1-d zone. You need to ensure high availability and replicate the server to the europe-west2-c zone using the

fewest steps possible. What should you

- A. Create a snapshot from the disk. Create a disk from the snapshot in the europe-west2-c zone. Create a new VM with that disk. 
- B. Create a snapshot from the disk. Create a disk from the snapshot in the europe-west1-d zone and then move the disk to europe-west2-c. Create a new VM with that disk.
- C. Use "gcloud" to copy the disk to the europe-west2-c zone. Create a new VM with that disk.
- D. Use "gcloud compute instances move" with parameter "--destination-zone europe-west2-c" to move the instance to the new zone.

Explanation:

Correct answer A

A is correct because this makes sure the VM gets replicated in the new zone.

B is not correct because this takes more steps than A.

C is not correct because this will generate an error, because gcloud cannot copy disks.

D is not correct because the original VM will be moved, not replicated.

Ask our Experts

Rate this Question?  

[View Queries](#)

open 

Question 10

Unattempted

Domain :Planning and configuring a cloud solution

You are working for a fast growing startup which has user base in US and Europe. Until now the company had all servers in Oregon region and has now decided to launch a replica of the entire infrastructure in EU region as well.

You started migration process with the least used server by creating snapshot of its disk and moving it to London region. Once the snapshot was copied you launched a fresh VM using the snapshot but discovered that few of the files were corrupted. What can be the reason?

- A. You cannot launch an instance from a snapshot.
- B. Snapshot was taken on a running VM. 

- C. As the region was changed all the files were not decrypted by properly which has resulted in data corruption.
- D. None of the above.

Explanation:

Answer B

Option B is correct: It is never recommended by Google to create a snapshot of running because this might corrupt data. You must always stop the instance and take a snapshot.

Option A is incorrect: You can definitely launch a VM from an existing snapshot.

Option C is incorrect: As the encryption was done using Google managed key, it is the responsibility of Google to decrypt and re-encrypt the files.

Option D is incorrect: As B is the correct option this is an invalid choice.

Link:

<https://cloud.google.com/compute/docs/disks/create-snapshots>

[Ask our Experts](#)[Rate this Question?](#)  [View Queries](#)[open ▾](#)**Question 11****Unattempted****Domain :Planning and configuring a cloud solution**

Being a Senior Cloud Engineer of the company, you are asked to launch a managed MySQL DB using custom VPC with network range of 172.17.0.0/16 on Google Cloud keeping security at utmost priority allowing team members to access it only over private channel.

You completed the setup with all the requirements, but developers are unable to access DB over the private network via application hosted on VM. What can be the reason?

- A. You did not whitelist 172.17.0.0/16 in the firewall. 
- B. This is a reserved IP range hence cannot be used.
- C. You need to use proxy to connect over private IP.
- D. None of the above.

Explanation:

Answer A

Option B is incorrect: The IP range 172.17.0.0/16 is usually used for docker bridge network. But it is not reserved and hence this is incorrect answer.

Option A is correct: You did not whitelist 172.17.0.0/16 in the firewall.

Option C is incorrect: You do not need proxy in this case because you are trying to access DB on the same network.

Option D is incorrect: As B is the correct option this is an invalid choice.

Link:

<https://cloud.google.com/sql/docs/mysql/configure-private-ip>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

[open ▾](#)

Question 12

Unattempted

Domain :Planning and configuring a cloud solution

A new client has approached your company with a requirement that they want to host a serverless on Google Cloud, but before that want you to give them price estimate of running serverless application using Cloud Function. Which all parameters will you consider while calculating price using Google Pricing Calculator?

- A. Type, Execution Time and Invocations.
- B. Memory, CPU, Execution Time and Invocations.
- C. Type, Bandwidth, Execution Time and Invocations. 
- D. Memory, Execution Time and Bandwidth.

Explanation:

Answer C

Option C is correct: You will need four parameters to give your client an estimate of running Cloud Functions. These are: Type, Bandwidth, Execution Time and Invocations per month.

Option A, B and D are incorrect because they are not complete set of parameters required to calculate estimate of running Cloud Functions on Google Cloud.

Link:

<https://cloud.google.com/functions/pricing>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 13

Unattempted

Domain :Planning and configuring a cloud solution

Your managers want you to suggest a low-cost storage option that could be used to share read-only data across multiple instances with high performance and no edge caching. Which storage option would you suggest?

- A. Cloud Storage 
- B. Persistent Disk
- C. Cloud Filestore
- D. Firebase

Explanation:

Answer A

B is incorrect: Persistent disks are used for read/write data across single instance with high performance and no edge caching.

A is correct: Cloud Storage is reliable object storage with an Option for no edge-caching service.

C is incorrect: Cloud Filestore is fully managed file storage service with high cost

D is incorrect: Firebase is scalable storage used for mobile app

Link:

<https://cloud.google.com/products/storage/>

[Ask our Experts](#)Rate this Question?  [View Queries](#)[open ▾](#)**Question 14****Unattempted**

Domain :Setting up a cloud solution environment

Your teammate launched 3 instances using gcloud compute instances create command with all the required flags. After few mins, you checked the console and found 0 instances in the GCE virtual machine section. How would you identify the project against which the command executed?

- A. gcloud auth application-default
- B. gcloud projects list
- C. gcloud debug
- D. gcloud config list 

Explanation:**Answer D**

D is correct: This command would list the properties for the currently active configurations.

B in incorrect: The following command will list all the projects that are created in a particular GCP account

A,C are incorrect: The following commands doesn't support the desired result.

Link:<https://cloud.google.com/sdk/gcloud/reference/config/list>[Ask our Experts](#)Rate this Question?  [View Queries](#)[open ▾](#)**Question 15****Unattempted**

Domain :Setting up a cloud solution environment

As a cloud engineer, you have been asked to upgrade the free trial of your account and rename it to a production-inventory-system. You are getting permission denied error while making the changes. Which of the following permissions will solve the problem?

- A. **billing.accounts.update** 
- B. **Billing.account.upgrade**
- C. **billing.account.update**
- D. **billing.accounts.upgrade**

Explanation:

Answer A

A is correct: The required permission is `billing.accounts.update` on Billing Account resource.

B, C and D are invalid choices or commands

Link:

<https://cloud.google.com/billing/docs/how-to/custom-roles>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 16

Unattempted

Domain :Planning and configuring a cloud solution

Your client wants to migrate an application to Google Cloud which has 15 TB of relational data. The database is growing rapidly by 10 GB everyday. In addition, to support the traffic, at least 10 read replicas are required. Which of the following service would you meet the requirements?

- A. **Cloud Storage**
- B. **BigQuery**
- C. **Cloud Spanner** 
- D. **Cloud Bigtable**

Explanation:**Answer C**

C is correct: Cloud Spanner is globally scalable, fully managed, enterprise relational database with automatic replication.

A is incorrect: Cloud Storage is reliable object storage with global edge-caching service

B is incorrect: BigQuery is a serverless cloud data warehouse

D is incorrect: BigTable is a petabyte scale, fully managed NoSQL database

Link:

<https://cloud.google.com/spanner>

[Ask our Experts](#)**Rate this Question?**  [View Queries](#)[open ▾](#)**Question 17****Unattempted****Domain :Planning and configuring a cloud solution**

Your client wants to develop a new cost effective web application that runs on serverless platform using Cloud Function, Cloud Storage, Pub/Sub and Cloud CDN. The expected data would be 20 GB. Which of the following database would be the most suitable schemaless option to support the serverless functionality?

- A. **Cloud BigTable**
- B. **BigQuery**
- C. **Cloud Spanner**
- D. **Cloud Firestore** 

Explanation:**Answer D**

D is correct: Cloud Firestore is a fast, fully managed, serverless, cloud-native NoSQL document database that simplifies storing, syncing, and querying data for your mobile, web, and IoT apps at global scale

A is incorrect: BigTable is a petabyte-scale, fully managed NoSQL database.

B is incorrect: BigQuery is serverless, highly scalable, cost effective data warehouse.

C is incorrect: Spanner is a fully managed relational database.

Link:

<https://cloud.google.com/firestore>

Ask our Experts

Rate this Question?  

View Queries

[open ▾](#)

Question 18

Unattempted

Domain :Deploying and implementing a cloud solution

You have created a pre-populated PersistentVolume disk as ReadOnlyMany, when you try to mount the volume to a POD, you get failed to mount error, what could be the most likely reason for this failure?

- A. You created a PersistentVolume but did not create a PersistentVolumeClaim 
- B. You have not specified the correct mount point in your pod.yaml file
- C. You entered an incorrect PersistentVolume name in the pod.yaml file
- D. The pod.yaml file doesn't have permission to mount the volume as it is created with ReadOnlyMany access.

Explanation:

Answer A

A is correct: In order to mount a persistent disk, you need to create a PersistentVolumeClaim after creating a PersistentVolume and then attach the PersistentVolumeClaim to the pod.

B is incorrect: A mount point is a user defined attribute and could be anything

C is incorrect: You cannot give a PersistentVolume name in pod.yaml, it should always be PersistentVolumeClaim

D is incorrect: ReadOnlyMany permissions does not derive the pod.yaml

Link:

<https://cloud.google.com/kubernetes-engine/docs/how-to/persistent-volumes/readonlymany-disks>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 19

Unattempted

Domain :Deploying and implementing a cloud solution

Which of the following command could be used to autoscale a replication controller "foo" with number of pods between 1 to 10 having target cpu utilization at 80%?

- A. `kubectl autoscale deployment foo --min=1 --max=10 --cpu-percent=80`
- B. `kubectl rc --autoscale foo --min=1 --max=10 --cpu-percentage=80`
- C. `kubectl deployment --autoscale foo --min=1 --max=10 --cpu-percent=80`
- D. `kubectl autoscale rc foo --max=10 --cpu-percent=80` 

Explanation:

Answer D

D is correct: This command is used to autoscale a replication controller with max 10 pods. It is important to specify --max flag as it defines the target pods that could be launched when cpu reaches 80%.

A is incorrect: The following command is used to autoscale a deployment and not rc

B & C is incorrect: Both the commands are incorrect.

Link:

<https://cloud.google.com/kubernetes-engine/docs/how-to/scaling-apps>

<https://kubernetes.io/docs/reference/generated/kubectl/kubectl-commands>

Ask our Experts

Rate this Question?  

[View Queries](#)[open ▾](#)**Question 20****Unattempted****Domain :Ensuring successful operation of a cloud solution**

Your client has a web application that is currently using a single GCE to read the messages from a pub/sub topic, process them and store them in Cloud Bigtable by its unique order id. When a user uploads an order to Cloud Storage, it triggers a Cloud Function that inserts a message to pub/subtopic. The traffic is expected to increase next week and as a cloud engineer, you are expected to resolve any possible auto scaling problem.

Which service is not likely to auto scale when the traffic is increased?

- A. Cloud Storage
- B. Cloud Bigtable
- C. GCE
- D. Pub/Sub

Explanation:**Answer C**

B. GCE is not a managed service and hence auto scaling is set up manually by using Managed Instance Group while architecting the infrastructure.

A, B and D are incorrect: These are the managed services and hence auto scaling is managed by GCP

Link:<https://cloud.google.com/compute/docs/autoscaler>[Ask our Experts](#)[Rate this Question? !\[\]\(40ea97637c564e3be05a46f598703620_img.jpg\) !\[\]\(1d6dc030007d6a92327cebda919704b6_img.jpg\)](#)[View Queries](#)[open ▾](#)**Question 21****Unattempted****Domain :Ensuring successful operation of a cloud solution**

You as a Senior Cloud Engineer doing POC on a business-critical application having database

deployed on a GCE virtual machine. The requirement is to have a disk that supports up to 35,000 Read IOPS per instance. Which of the following disks is suitable to meet the requirement?

- A. Standard Persistent Disk
- B. SSD Persistent Disk 
- C. Local SSD (SCSI)
- D. Local SSD (NVMe)

Explanation:

Answer B

B is correct: SSD Persistent Disk supports up to 15000-60000 IOPS per instance

A is incorrect: Standard Persistent Disk supports up to 3000 IOPS

C is incorrect: SCSI supports up to 400000 IOPS

D is incorrect: NVMe supports up to 680000 IOPS

Link:

<https://cloud.google.com/compute/docs/disks>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 22

Unattempted

Domain :Configure access and security

John & co has deployed an application using Google App Engine standard environment. You have been asked to update the cron schedules and default cookie expiration time, which of the following predefined role has access to update default cookie expiration but no access to update cron schedules?

- A. App Engine Service Admin
- B. App Engine Admin 

- C. App Engine Deployer
- D. App Engine Code Viewer

Explanation:

Answer B

B is correct: As per google docs, App Engine Admin has access to update default cookie expiration but no access to update cron schedules

A, C & D are incorrect: None of the other predefined roles have permission to perform either of the operations

Link:

<https://cloud.google.com/appengine/docs/standard/python/roles>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 23

Unattempted

Domain :Configure access and security

You company has uploaded some business critical documents to Cloud Storage and your project manager wants you to restrict access to the objects by using ACLs. Which of the following permission would allow you to update the object ACLs?

- A. storage.objects.update
- B. storage.objects.setIamPolicy 
- C. storage.objects.create
- D. storage.objects.getIamPolicy

Explanation:

Answer B

B is correct: As per google docs, storage.objects.setIamPolicy allows user to update object ACLs.

A is incorrect: Update object metadata, excluding ACLs

C is incorrect: Add new objects to the bucket

D is incorrect: Read object ACLs, returned as IAM policies

Link:

<https://cloud.google.com/storage/docs/access-control/iam-permissions>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 24

Unattempted

Domain :Setting up a cloud solution environment

Which of the following role provide granular access for a specific service and is managed by GCP?

- A. Custom
- B. Predefined
- C. Admin
- D. Primitive

Explanation:

Answer B

B is correct: Predefined roles are managed roles and provide service specific access.

A is incorrect: Provide granular access for a specific service but managed by users

C is incorrect: Invalid role type

D is incorrect: Concrete roles, existed prior to IAM

Link:

<https://cloud.google.com/iam/docs/understanding-roles>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 25**Unattempted**

Domain :Setting up a cloud solution environment

When working on billing dashboard, which of the following option would lead you to cost details for a specific invoice month?

- A. Report
- B. Cost Breakdown
- C. Cost Table 
- D. Billing Export

Explanation:

Answer C

C is correct: The cost table gives a detailed tabular view of monthly costs for a given invoice.

A is incorrect: Lets you view your costs at a glance to discover and analyze trends.

B is incorrect: Gives you a waterfall view of your costs

D is incorrect: Billing exports is used for monitoring, analyzing and optimizing costs.

Link:

<https://cloud.google.com/billing/docs/how-to/cost-table>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 26**Unattempted**

Domain :Planning and configuring a cloud solution

Your client wants to serve content using Cloud Storage that allow similar performance advantage as regions, but also wants higher availability that comes with being geo-redundant. You plan to use dual region to meet his requirement. Choose the regions that does not fall under dual region compatibility. (Select 2)

- A. nam4
- B. us
- C. eur4
- D. eu

Explanation:

Answer B and D

B and D correct: us and eu are supported under multi-region and not in dual region

A and C are incorrect: The following 2 regions are currently supported in dual region.

Link:

<https://cloud.google.com/storage/docs/locations>

[Ask our Experts](#)[Rate this Question?](#)  [View Queries](#)[open](#) ▾**Question 27****Unattempted**

Domain :Planning and configuring a cloud solution

How many maximum instance-hours are supported by an App Engine as free daily usage quota when scaling an instance using Automatic Scaling type?

- A. 28
- B. 8
- C. 35
- D. 10

Explanation:

Answer A

A is correct: As per google documentation, automatic scaling has a free daily usage quota of 28 instance hours

B is incorrect: Free usage quota of 8 instance-hours is supported by Manual and Basic scaling

C and D are incorrect: Invalid choices

Link: <https://cloud.google.com/appengine/quotas#Instances>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 28

Unattempted

Domain :Deploying and implementing a cloud solution

You created a VPC with a CIDR block of 10.40.0.0/16 with 2 subnets of CIDR range 10.40.1.0/24 and 10.40.2.0/24. What would be the default routes within this virtual private connection having the broadest CIDR range?

- A. 10.40.0.0/0
- B. 10.40.3.0/24
- C. 0.0.0.0/0 
- D. 0.0.0.0/16

Explanation:

Answer C

C is correct: The system generated default route is 0.0.0.0/0 as it is the broadest possible range.

A, B and D are incorrect routes

Link:

<https://cloud.google.com/vpc/docs/routes>

Ask our Experts

Rate this Question?  

[View Queries](#)

[open ▾](#)

Question 29

Unattempted

Domain :Deploying and implementing a cloud solution

There are 5 VPC networks in your staging project created for 5 different applications. Each network has its own CIDR range and firewall rules. You are asked to list the firewall rules of network 3 to perform an update. Select the appropriate response.

- A. `gcloud compute firewall-rules list --filter network=network 3` 
- B. `gcloud vpc network=network 3 --list firewall-rules`
- C. `gcloud compute network=network 3 --list firewall-rules`
- D. `gcloud vpc firewall-rules list --filter network=network 3`

Explanation:

Answer A

A is correct: The firewall-rules are defined under VPC but when using gcloud, they can only be accessed by calling compute api and this the following command lists the firewall-rules for the given network.

B, C & D is incorrect: These commands are invalid.

Link:

<https://cloud.google.com/vpc/docs/using-firewalls>

Ask our Experts

Rate this Question?  

[View Queries](#)

[open ▾](#)

Question 30

Unattempted

Domain :Ensuring successful operation of a cloud solution

Your application servers are deployed in private subnet and you connect to an application instance using a bastion host, you connect to the bastion host using Console Shell and then ssh to the respective application instance. Unfortunately, Cloud Shell is taking a lot of time to respond and you manager wants you to ssh to the instance to fix an application specific configuration. What command would you use to SSH into the server from public end?

- A. gcloud connect ssh username@application-instance
- B. gcloud compute ssh username@bastion-host 
- C. gsutil connect ssh username@bastion-host
- D. gsutil compute ssh username@application-server

Explanation:

Answer B

B is correct. The correct syntax is gcloud compute ssh + [USERNAME]+@[SERVERNAME]. From public end, you can only connect to bastion host, hence you need to connect to bastion host first and then connect to application instance.

A, C and D are incorrect: Incorrect commands

Link:

<https://cloud.google.com/compute/docs/instances/connecting-to-instance>

[Ask our Experts](#)[Rate this Question? !\[\]\(8197433878765d452af236394c75d433_img.jpg\) !\[\]\(1bee590b91307765d8b3434c46ed84a8_img.jpg\)](#)

[View Queries](#)[open !\[\]\(663232dec2f5f9f4a498ccdac04c36f6_img.jpg\)](#)

Question 31**Unattempted**

Domain :Configure access and security

You deployed 10 micro services using Google Kubernetes Engine. The command kubectl run deployed the micro services in different pods, 2 week later your manager asked you to delete the pods as the services were no longer needed. Which of the following options is the best way to delete the pods?

- A. gcloud container clusters delete
- B. gcloud container clusters delete pod

C. `kubectl delete deployments/<name> services/<name>`

D. `kubectl remove pods`

Explanation:

Answer C

C is correct: The best way is to delete the deployment is deleting the deployment and the services.

A is incorrect: This will delete the entire cluster and all the deployments within it

B and D are incorrect: Incorrect Commands

Link:

<https://cloud.google.com/kubernetes-engine/docs/quickstart>

[Ask our Experts](#)

Rate this Question?

[View Queries](#)

[open](#) ▾

Question 32

Unattempted

Domain :Configure access and security

ASM Info has deployed an application using Google App Engine flexible environment and you have been asked to ssh the VM instance in flexible environment, update the app configuration and if required, enable and disable the App Engine application, which of the following predefined role has access to perform these tasks?

A. App Engine Admin

B. App Engine Service Admin

C. App Engine Deployer

D. App Engine Code Viewer

Explanation:

Answer A

A is correct: As per google docs, App Engine Admin has access to perform these tasks

B, C & D are incorrect: None of the other predefined roles have permission to perform either of the operations

Link:

<https://cloud.google.com/appengine/docs/standard/python/roles>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 33

Unattempted

Domain :Configure access and security

You company has developed an application to accelerate the efficiency in making sales calls about the product of your client. Everytime a recipient is called, log sink is used to export the record to a destination. Which of the following destinations are supported? Choose any 2.

- A. BigQuery 
- B. Cloud Spanner
- C. Pub/Sub 
- D. Cloud Bigtable

Explanation:

Answer A and C

A and C are correct: Log sinks can be exported to Cloud Storage, Pub/Sub and BigQuery only.

B and D are incorrect: Log sinks cannot be exported to either of the two

Link:

<https://cloud.google.com/logging/docs/export/>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 34**Unattempted****Domain :Setting up a cloud solution environment**

You created a bucket in cloud storage and uploaded some files and then enabled object versioning on it. The files you have already added will have which of the following version?

- A. 1
- B. Null 
- C. 0
- D. -1

Explanation:**Answer B**

B is correct: By default, the object will have null version as versioning was enabled after uploading the files.

A is incorrect: The new object uploaded will have metageneration number 1

C and D are incorrect: Invalid choices

Link:<https://cloud.google.com/storage/docs/object-versioning>[Ask our Experts](#)**Rate this Question?**  [View Queries](#)[open ▾](#)**Question 35****Unattempted****Domain :Setting up a cloud solution environment**

Your company has 5 TB of testing data stored in the production database of a testing tool name Quality Center. The data is being used to create a real time analytics system which is causing slow response to the tester while using the tool. What should you do to improve the load on the database?

- A. Set up Multi AZ

- B. Set up a read replica 
- C. Scale the database instance
- D. Run the analytics query only weekends

Explanation:

Answer: B

B is correct: Read replica can be used to run all the queries related to Analytics system.

A is incorrect: Multi AZ would improve the availability of the data

C is incorrect: This will not add any value

D is incorrect: The analytics will not be realtime.

Link:

<https://cloud.google.com/sql/docs/mysqlreplication/>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

[open](#) 

Question 36

Unattempted

Domain :Planning and configuring a cloud solution

Your inventory application has 2 backend api servers launched using Auto Scaling Groups, you have been asked to load balance UDP, TCP and SSL traffic on ports that are not supported by the TCP proxy and SSL proxy load balancer. Which of the following Load Balancer would you use?

- A. HTTP(S) Load Balancer
- B. SSL Proxy Load Balancer
- C. TCP Proxy Load Balancer
- D. Network Load Balancer 

Explanation:

Answer: D

D correct: You can use Network Load Balancing to load balance UDP, TCP, and SSL traffic on ports that are not supported by the TCP proxy and SSL proxy load balancers.

A is incorrect: An HTTPS load balancer uses a target HTTPS proxy instead of a target HTTP proxy

B is incorrect: With SSL Proxy Load Balancing, SSL connections are terminated at the load balancing layer then proxied to the closest available instance group.

C is incorrect: Cloud TCP Proxy Load Balancing is intended for non-HTTP traffic

Link:

<https://cloud.google.com/load-balancing/docs/network/>

Ask our Experts

Rate this Question?  

[View Queries](#)

open ▾

Question 37

Unattempted

Domain :Planning and configuring a cloud solution

You are distributing the traffic between a fleet of VMs with in your VPC using Internal TCP/UDP Load Balancer. Which of the following specification does not support the selected Load Balancing Type?

- A. Preserved Client IP
- B. Global Availability 
- C. Internal Load Balancing
- D. Any Destination Ports

Explanation:

Answer: B

B is correct: Internal TCP/UDP Load Balancer are available within a Region and not Globally

A is incorrect: Internal TCP/UDP Load Balancer preserves client IP address

C is incorrect: Internal TCP/UDP Load Balancer are internally faced

D is incorrect: Internal TCP/UDP Load Balancer allow access to any destination port

Link:

<https://cloud.google.com/load-balancing/docs/choosing-load-balancer>

Ask our Experts

Rate this Question?  

[View Queries](#)

[open ▾](#)

Question 38

Unattempted

Domain :Planning and configuring a cloud solution

You updated the metadata of an object stored in Reports-PDF bucket using google cloud SDK. You want to ensure that the metadata has been updated for that object before you confirm the changes to the development team. Which of the following google cloud SDK command would you use?

- A. `gsutil stat gs://Reports-PDF/[OBJECT_NAME]` 
- B. `gsutil getmeta gs://Reports-PDF/[OBJECT_NAME]`
- C. `gsutil setmeta gs://Reports-PDF/[OBJECT_NAME]`
- D. `gsutil meta gs://Reports-PDF/[OBJECT_NAME]`

Explanation:

Answer: A

A is correct: The stat command will output details about the specified object URLs.

B is incorrect: getmeta is not a valid GCP command

C is incorrect: setmeta is used to update the metadata of an object with additional flags

D is incorrect: meta is not a valid GCP command

Link:

<https://cloud.google.com/storage/docs/viewing-editing-metadata>

Ask our Experts

Rate this Question?  

[View Queries](#)[open ▾](#)**Question 39****Unattempted****Domain :Planning and configuring a cloud solution**

Which of the following command could be used to change the storage class of an object in Cloud Storage?

- A. gsutil rewrite -c [STORAGE_CLASS] gs://Reports-PDF/[OBJECT_NAME]
- B. gsutil rewrite -s [STORAGE_CLASS] gs://Reports-PDF/[OBJECT_NAME] 
- C. gsutil rewrite -l [STORAGE_CLASS] gs://Reports-PDF/[OBJECT_NAME]
- D. gsutil rewrite -k [STORAGE_CLASS] gs://Reports-PDF/[OBJECT_NAME]

Explanation:**Answer: B**

B is correct: The gsutil rewrite command rewrites cloud objects, applying the specified transformations to them. The transformation(s) are atomic and applied based on the input transformation flags (-s). Object metadata values are preserved unless altered by a transformation.

A is incorrect: Invalid flag and command.

C is incorrect: Causes gsutil to read the list of objects to rewrite from stdin.

D is incorrect: Rewrite objects with the current encryption key specified in your boto configuration file

Link:<https://cloud.google.com/storage/docs/gsutil/commands/rewrite>[Ask our Experts](#)[Rate this Question? !\[\]\(aca0204fa99dd7cfeb7f9c20d47b755a_img.jpg\) !\[\]\(32bf2efd45946a576bcdc55cab7ad26e_img.jpg\)](#)[View Queries](#)[open ▾](#)**Question 40****Unattempted****Domain :Deploying and implementing a cloud solution**

Among the list of permissions attached below, which of the following permissions are required to manage SSH keys on the project while setting project-wide metadata to access the instance if the OS Login is not working? Select 2.

- A. `compute.instance.setMetadata`
- B. `compute.project.setCommonInstanceMetadata` 
- C. `iam.serviceAccounts.actAs` 
- D. `iam.setsshMetadata.allInstances`

Explanation:

Answer: B and C

B and C are correct: The following 2 permissions are required on the project if setting project-wide metadata.

A is incorrect: Permission is required on the instance if setting metadata on the instance.

D is incorrect permission.

Link:

<https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 41

Unattempted

Domain :Deploying and implementing a cloud solution

You have been assigned to a new health application project where the backend instances are deployed using Managed Instance Groups. There are 4 instances running. The MIG is not set to automatically scale and you are asked to resize a group to handle an expected increase in traffic. Which of the following statements about MIG is incorrect?

- A. If you increase the size, the managed instance group uses the current instance template to add new instances.

- B. The group deletes instances with a current action of DELETING, CREATING, and RECREATING before it deletes instances that are running with no scheduled actions.
- C. If the group is part of a backend service that has enabled connection draining, it can take up to 60 seconds after the connection draining duration has elapsed before the VM instance is removed or deleted.
- D. When updating a MIG, no more than 500 instances can be specified in a single request. 

Explanation:

Answer: D

D is correct: When updating a MIG, up to 1000 instances can be specified in a single request.

A, B & C is incorrect: These statements are correct with respect to MIG but incorrect w.r.t to the question.

Link:

https://cloud.google.com/compute/docs/instance-groups/creating-groups-of-managed-instances#resize_managed_group

[Ask our Experts](#)[Rate this Question?](#)  [View Queries](#)[open ▾](#)**Question 42****Unattempted**

Domain :Deploying and implementing a cloud solution

Your company has an application that is deployed using serverless architecture by making use of Cloud Function as backend code, Pub/Sub, Endpoints and serve the static content via Cloud Storage. Your application is used heavily and you were informed about an issue with respect to the Cloud Function. You realised that the issue is because of invocation limit per second. What is the default limit set by GCP to invoke a function per second?

- A. 65
- B. 1000 
- C. 80
- D. 70

Explanation:

Answer: B

B is correct: The default invocation limit set by GCP is 1000 per seconds

A, C and D are incorrect: Invalid Limit options

Link: https://cloud.google.com/functions/quotas#additional_quotas_for_background_functions

[Ask our Experts](#)[Rate this Question? !\[\]\(23c95a79ea2fdecb85c88c5fd848f31c_img.jpg\) !\[\]\(a2d4ede976271f5b3eeb874ceb1d5a35_img.jpg\)](#)[View Queries](#)[open !\[\]\(110cc478cf3551d7691515804cdd34ca_img.jpg\)](#)**Question 43****Unattempted****Domain :Deploying and implementing a cloud solution**

While migrating traffic gradually between the 2 versions of your Flexible App Engine environment, you have encountered a spike in latency for loading request. Which of the following statement could be a cause of this issue?

- A. The warm up requests are disabled, hence the traffic migration is not working as expected

- B. Gradual traffic migration between the versions running in flexible environment is not supported. 
- C. Deploying a new version with the same name as an existing version is causing the spike in latency.
- D. You can migrate the traffic gradually only for the versions that are deployed in different flexible environments.

Explanation:

Answer: B

B is correct: Gradual traffic migration is not supported in flexible environment.

A, C and D are incorrect: Choices are valid only for immediate traffic migration.

Link:

https://cloud.google.com/appengine/docs/standard/python/migrating-traffic#migrating_traffic_to_a_new_version

[Ask our Experts](#)[Rate this Question?](#)  [View Queries](#)[open ▾](#)**Question 44****Unattempted**

Domain :Ensuring successful operation of a cloud solution

A health care company that provides medical service to the users want to track their network forensics, real-time security analysis and optimize the expense. The manager would like to track the network sent from and received by VM instances. What do you suggest they do?

- A. Pub/Sub
- B. VPC Flow Logs 
- C. VPC Monitoring
- D. Cloud Storage

Explanation:

Answer: B

B is correct. Flow Logs are used to track network related findings.

A is incorrect: Pub/Sub is used to stream log data to other services

B is incorrect: VPC Monitoring is an invalid option

D is incorrect: Cloud Storage is used to push/storage objects.

Link:

<https://cloud.google.com/vpc/docs/using-flow-logs>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 45

Unattempted

Domain :Ensuring successful operation of a cloud solution

Your client has prepared a new company policy in which each developer must sign a Contributor License Agreement (CLA) before code changes are committed to any version control repository. You have been asked to check each commit in a repository that includes the policy and your manager has also provided you with node.js code. Which of the following services can help you implement this solution?

- A. Task Scheduler
- B. Cloud Function 
- C. Pub/Sub
- D. GKE

Explanation:

Answer: B

B is correct: Cloud Function can be used to retrieve commits, analyze code, committers and perform creative tasks such as checking a CLA.

A is incorrect: Task Scheduler is the scheduled task itself. Beneficial to perform repetitive tasks.

C is incorrect: It is used only to publish and subscribe messages

D is incorrect: Container services are microservices that require additional configuration efforts.

Link:

<https://cloud.google.com/functions/>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 46

Unattempted

Domain :Ensuring successful operation of a cloud solution

A development team works with two Cloud Functions using node.js code. Each function corresponds to environments for development and production. The code is same except for the Cloud SQL database values used in each function. Team wants to maintain code in a clean and reusable fashion and decides to pass the database value during function execution. Which feature will allow you to do this?

- A. **Timeouts**
- B. **Service Accounts**
- C. **Environment Variables** 
- D. **Metadata**

Explanation:

Answer: C

C is correct: Environment variables for cloud functions enable you to dynamically pass settings to your function code and libraries, without making changes to your code

A is incorrect: Is the duration for which the function will execute.

B is incorrect: Allows you to configure appropriate permissions

D is incorrect: Metadata is an invalid option wrt to Cloud Function

Link:

<https://cloud.google.com/functions/docs/env-var>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 47**Unattempted**

Domain :Ensuring successful operation of a cloud solution

What does the CIDR 10.0.2.0/26 correspond to?

- A. 10.0.2.0 - 10.0.2.26
- B. 10.0.2.0 - 10.0.2.63 
- C. 10.0.0.0 - 10.0.63.0
- D. 10.0.2.0 - 10.0.0.26

Explanation:

Answer: B

B is correct: /26 means 64 IPs ($=2^{(32-26)} = 2^6$), means only the last digit can change.

A, C and D are incorrect CIDR choices.

Link:

<https://cloud.google.com/vpc/docs/vpc>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 48**Unattempted**

Domain :Configure access and security

Your client hosts a static website on Cloud Storage written in HTML, CSS, JavaScript. The site targets users in North America, the usage of website has grown worldwide and hundreds of thousands of visitors access it monthly. Visitors from different parts of the world are experiencing slow performance due to latency while users in the United States experience normal response times. What service can mitigate this issue?

- A. Load Balancer
- B. Cloud DNS
- C. Cloud CDN 
- D. Storage Caching

Explanation:

Answer: C

C is correct: Google Cloud CDN leverages Google's globally distributed edge points of presence to accelerate content delivery for websites and applications served out of Google Compute Engine and Google Cloud Storage.

A is incorrect: Distributed loads between 2 or more instances

B is incorrect: Google Cloud DNS is a scalable, reliable, and managed authoritative Domain Name System (DNS) service running on the same infrastructure as Google.

D is incorrect: The service doesn't offer Caching.

Link:

<https://cloud.google.com/storage/docs/hosting-static-website>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 49

Unattempted

Domain :Configure access and security

Your company's data center has a CIDR of 10.20.0.0/10 and your remote office has 172.50.0.0/16. You have been asked to plan a migration of the workloads from your company's data center to GCP. Which of the following CIDR range would you select for your VPC in GCP?

- A. **172.50.0.0/12**
- B. **10.20.0.0/16**
- C. **10.40.0.0/16**
- D. **172.40.4.0/18** 

Explanation:

Answer: D

The company's data center CIDR 10.20.0.0/10 will have the following IP range:

10.0.0.1 - 10.63.255.254

And the above would overlap the subnet in letter C: 10.40.0.0/16, which has the following IP range:

10.40.0.1 - 10.40.255.254

Hence, the correct answer should be D: 172.40.4.0/18, which has the following IP range:

170.40.0.1 - 170.40.63.254

Link:

<https://cloud.google.com/vpc/docs/configure-alias-ip-ranges>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▼

Question 50

Correct

Domain :Planning and configuring a cloud solution

You have been hired as a contractor by one of the travel technology company who is planning to containerize their existing applications in such a way that they can perform a lift and shift very easily in future if they plan to move away from Google Cloud. Which service will best suit this case?

- A. **Cloud Function**
- B. **App Engine Standard**
- ✓ C. **Kubernetes Engine** 
- D. **CloudRun**

Explanation:

Answer C

Option C is correct: Kubernetes is an open-source solution supported by major cloud platforms and will be very easy for company to move out at later stage if required because they can use the same config files with very minor changes.

Option A is incorrect: With Cloud Function you will have to make changes to your function as per the requirement of vendor.

Option B is incorrect: Moving out of App Engine Standard will require you to use either Docker Swarm or Kubernetes for managing your containers.

Option D is incorrect: CloudRun is a service used to run stateless and serverless kube containers at scale. This could have been the option but as this service does not require you to write config files and when you plan to move out of GCP you will have to create them hence, consuming more time and efforts to migrate.

Link:

<https://cloud.google.com/kubernetes-engine/>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open 

[Finish Review](#)

Certification	Company	Support	Join us on Slack!
Cloud Certification	Become Our Instructor	Contact Us	 Join our open Slack community and get your queries answered instantly! Our experts are online to answer your questions!
Java Certification	Support	Help Topics	
PM Certification	Discussions		
Big Data Certification	Blog		Follow us
	Business		  

© Copyright 2021. Whizlabs Software Pvt. Ltd. All Right Reserved.